

Running Julia code in parallel with MPI


Lessons learned

Lars Christmann^{*†}, Niklas Neher^{*},
Michael Schlottke-Lakemper^{*}

JuliaCon 2022, 26th July 2022


^{*} High-Performance Computing Center Stuttgart (HLRS),
University of Stuttgart, Germany

[†] University of Cologne, Germany



Does it scale?

- Julia code for computational fluid dynamics: ✓ ([Trixi.jl](#))
- Parallelized with MPI: ✓
- Supercomputer available: ✓ ([Hawk @ HLRS](#))



©Trixi.jl authors, Ben Derzian, HLRS

What is Trixi.jl?

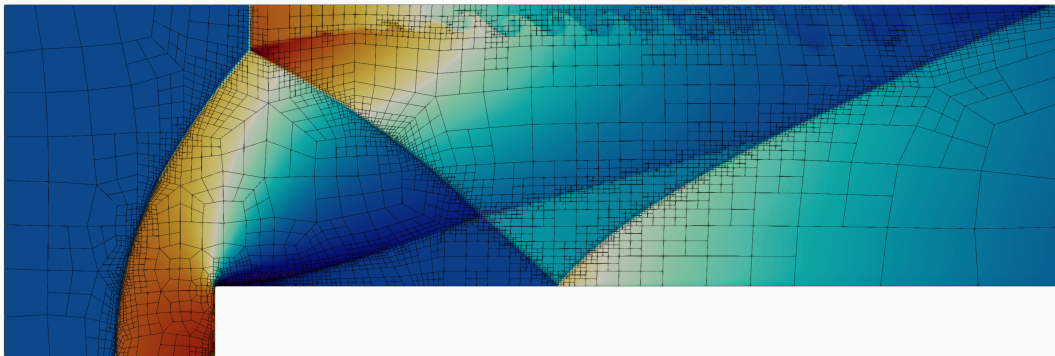
Trixi.jl

- **Adaptive** numerical simulation framework
- Focus on **high-order methods** for **computational fluid dynamics**
- **Hybrid parallelization** with MPI and multithreading (Polyester.jl)



<https://github.com/trixi-framework/Trixi.jl>

Trixi.jl simulation of front-facing step in wind tunnel at Mach 3



Credit: Andrew R. Winters, <https://www.youtube.com/watch?v=glAug1aIxio>

Hawk and Vulcan

- High-Performance Computing Center Stuttgart (HLRS)
- Hawk → capability computing (AMD EPYC Rome, ~720k cores)
- Vulcan → industry computing (various Intel Xeon, 10k-20k cores)
- Julia 1.7.2 available as modules

<https://www.hlr.de/solutions/systems>



HLRS Hawk



HLRS Vulcan

Where to put the Julia folder

- `JULIA_DEPOT_PATH` envvar specifies location for package files, precompilation files. . .
- Put it in `home directory` (NFS)?
 - Julia crashed when using >2000 MPI ranks
- Put it on `parallel file system` (Lustre)?
 - >10,000 small files
 - Issues: slow on Pkg update, quota limits



Lessons learned

Figure out depot location before doing parallel runs.

Observations for MPI and multithreading

- Generally **high memory usage** of Julia runtime: Trixi.jl + OrdinaryDiffEq.jl
 - 360 MiB after loading
 - 510 MiB after compilation
- Hybrid parallelization (MPI + multithreading) support may be **system dependent**
- Non-standard MPI implementations **might** require additional efforts

Lessons learned

Test which MPI implementation to use.

Verify that multithreading works as expected.

Using Julia with a system MPI may require user action

- Need to **replace** MPI-enabled JLL-provided binaries by **system binaries** (e.g., MPI.jl, HDF5.jl)
- Need to **regenerate C bindings** for libraries due to MPI ABI change (e.g., P4est.jl)

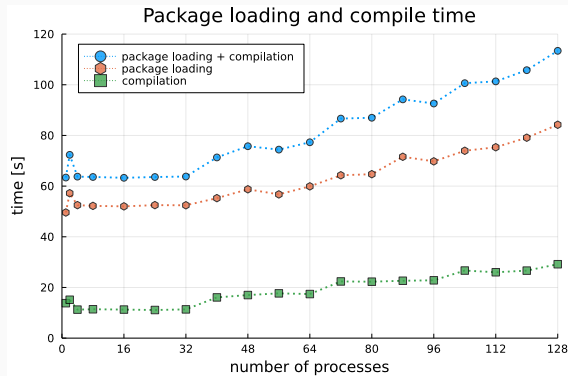
Lessons learned

Document setup procedures for your users and yourself.

Have a look at MPItrampoline (<https://github.com/eschnett/MPItrampoline>).

Runtime overhead due to code loading/compilation

- **Code loading** times increase due to file system pressure
- Also **code compilation** can be slower in parallel
- General rule: **non-parallel operations** increase computational cost



Lessons learned

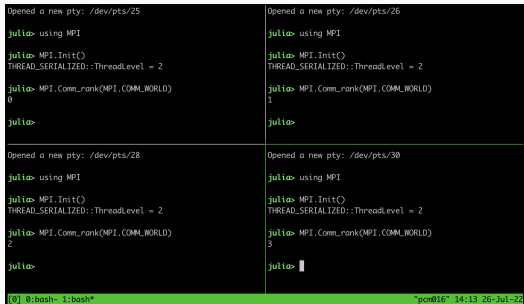
Use PackageCompiler.jl to create custom Julia sysimage

(<https://github.com/JuliaLang/PackageCompiler.jl>).

Try to use HPC/Spindle to avoid I/O bottleneck (<https://github.com/hpc/Spindle>).

Restrictions for interactive computing with MPI

- No inherent MPI support for the REPL
- Long load/compilation times
- Parallel development is sub-optimal (debugging is even less fun)



```
Opened a new pty: /dev/pts/25
julia> using MPI
julia> MPI.Init()
THREAD_SERIALIZED::ThreadLevel = 2
julia> MPI.Comm_rank(MPI.COMM_WORLD)
0
julia>

Opened a new pty: /dev/pts/26
julia> using MPI
julia> MPI.Init()
THREAD_SERIALIZED::ThreadLevel = 2
julia> MPI.Comm_rank(MPI.COMM_WORLD)
1
julia>

Opened a new pty: /dev/pts/28
julia> using MPI
julia> MPI.Init()
THREAD_SERIALIZED::ThreadLevel = 2
julia> MPI.Comm_rank(MPI.COMM_WORLD)
2
julia>

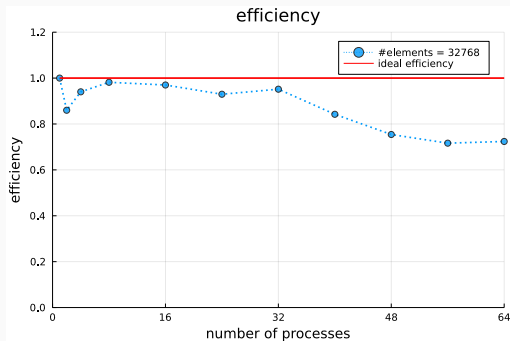

Opened a new pty: /dev/pts/30
julia> using MPI
julia> MPI.Init()
THREAD_SERIALIZED::ThreadLevel = 2
julia> MPI.Comm_rank(MPI.COMM_WORLD)
3
julia>

[0] 0:~ 1:~ bash* "pcn016" 14:13 26-Jul-22
```

Lessons learned

Run MPI in tmux via tmpi as a workaround (<https://github.com/Azrael3000/tmpi>).
Be patient.

Strong scaling results for Trixi.jl



Good scalability from 1 up to 32 MPI ranks, acceptable scalability until 64 MPI ranks

Lessons learned

Parallel performance of Julia with MPI behaves similarly to C++/Fortran codes.

Conclusions

- Running Julia code in parallel with MPI **works!**
- **Additional challenges** compared to traditional HPC languages
- Parallel development toolchain **not as mature yet**
- Is Julia with MPI the right tool for your HPC needs: **it depends**

Lessons learned

Interact with the great Julia for HPC community!

Thank you for your interest!